

```
In [9]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
df=pd.read_csv(r"C:/Users/DELL/\Downloads\used_cars_data.csv")
df
```

Out[9]:

	S.No.	Name	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	Owner_T
0	0	Maruti Wagon R LXI CNG	Mumbai	2010	72000	CNG	Manual	f
1	1	Hyundai Creta 1.6 CRDi SX Option	Pune	2015	41000	Diesel	Manual	f
2	2	Honda Jazz V	Chennai	2011	46000	Petrol	Manual	f
3	3	Maruti Ertiga VDI	Chennai	2012	87000	Diesel	Manual	f
4	4	Audi A4 New 2.0 TDI Multitronic	Coimbatore	2013	40670	Diesel	Automatic	Sec
...
7248	7248	Volkswagen Vento Diesel Trendline	Hyderabad	2011	89411	Diesel	Manual	f
7249	7249	Volkswagen Polo GT TSI	Mumbai	2015	59000	Petrol	Automatic	f
7250	7250	Nissan Micra Diesel XV	Kolkata	2012	28000	Diesel	Manual	f
7251	7251	Volkswagen Polo GT TSI	Pune	2013	52262	Petrol	Automatic	T
7252	7252	Mercedes-Benz E-Class 2009-2013 E 220 CDI Avan...	Kochi	2014	72443	Diesel	Automatic	f

7253 rows × 14 columns



In [10]: `df.head()`

Out[10]:

	S.No.	Name	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	Owner_Type
0	0	Maruti Wagon R LXI CNG	Mumbai	2010	72000	CNG	Manual	First
1	1	Hyundai Creta 1.6 CRDi SX Option	Pune	2015	41000	Diesel	Manual	First
2	2	Honda Jazz V	Chennai	2011	46000	Petrol	Manual	First
3	3	Maruti Ertiga VDI	Chennai	2012	87000	Diesel	Manual	First
4	4	Audi A4 New 2.0 TDI Multitronic	Coimbatore	2013	40670	Diesel	Automatic	Second

In [11]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7253 entries, 0 to 7252
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   S.No.                 7253 non-null  int64
1   Name                  7253 non-null  object
2   Location              7253 non-null  object
3   Year                  7253 non-null  int64
4   Kilometers_Driven     7253 non-null  int64
5   Fuel_Type             7253 non-null  object
6   Transmission          7253 non-null  object
7   Owner_Type            7253 non-null  object
8   Mileage               7251 non-null  object
9   Engine                7207 non-null  object
10  Power                 7207 non-null  object
11  Seats                 7200 non-null  float64
12  New_Price             1006 non-null  object
13  Price                 6019 non-null  float64
dtypes: float64(2), int64(3), object(9)
memory usage: 793.4+ KB
```

```
In [12]: df.describe()
```

```
Out[12]:
```

	S.No.	Year	Kilometers_Driven	Seats	Price
count	7253.000000	7253.000000	7.253000e+03	7200.000000	6019.000000
mean	3626.000000	2013.365366	5.869906e+04	5.279722	9.479468
std	2093.905084	3.254421	8.442772e+04	0.811660	11.187917
min	0.000000	1996.000000	1.710000e+02	0.000000	0.440000
25%	1813.000000	2011.000000	3.400000e+04	5.000000	3.500000
50%	3626.000000	2014.000000	5.341600e+04	5.000000	5.640000
75%	5439.000000	2016.000000	7.300000e+04	5.000000	9.950000
max	7252.000000	2019.000000	6.500000e+06	10.000000	160.000000

```
In [13]: df.isnull().sum()
```


```
Out[13]: S.No.          0
Name          0
Location      0
Year          0
Kilometers_Driven  0
Fuel_Type     0
Transmission  0
Owner_Type    0
Mileage       2
Engine        46
Power         46
Seats         53
New_Price     6247
Price        1234
dtype: int64
```

```
In [14]: df.fillna(value=0,inplace=True)
```

```
In [15]: df=df[['Year','Price']]
#Taking only the selected two attributes from the dataset
df.columns=['year','price']
#Renaming the columns for easier writing of the code
```

```
In [16]: sns.lmplot(x='year',y='price',data=df,order=2,ci=None)
```

```
Out[16]: <seaborn.axisgrid.FacetGrid at 0x22adce76990>
```




```
In [17]: X=np.array(df['year']).reshape(-1,1)
```

```
In [18]: y=np.array(df['price']).reshape(-1,1)
```

```
In [19]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25)
reg=LinearRegression()
reg.fit(X_train,y_train)
print(reg.score(X_test,y_test))
```


```
0.06439829841709799
```

```
In [20]: y_pred=reg.predict(X_test)
plt.scatter(X_test,y_test,color='b')
plt.plot(X_test,y_pred,color='k')
plt.show()
```



```
In [21]: df500=df[:][:500]  
sns.lmplot(x='year',y='price',data=df500,order=1,ci=None)
```


```
Out[21]: <seaborn.axisgrid.FacetGrid at 0x22adfa4f010>
```



```
In [22]: df500.fillna(method='ffill',inplace=True)
X=np.array(df500['year']).reshape(-1,1)
y=np.array(df500['price']).reshape(-1,1)
df500.dropna(inplace=True)
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25)
reg=LinearRegression()
reg.fit(X_train,y_train)
print("Regression:",reg.score(X_test,y_test))
y_pred=reg.predict(X_test)
plt.scatter(X_test,y_test,color='b')
plt.plot(X_test,y_pred,color='k')
plt.show
```

Regression: 0.17291193573932195

Out[22]: <function matplotlib.pyplot.show(close=None, block=None)>

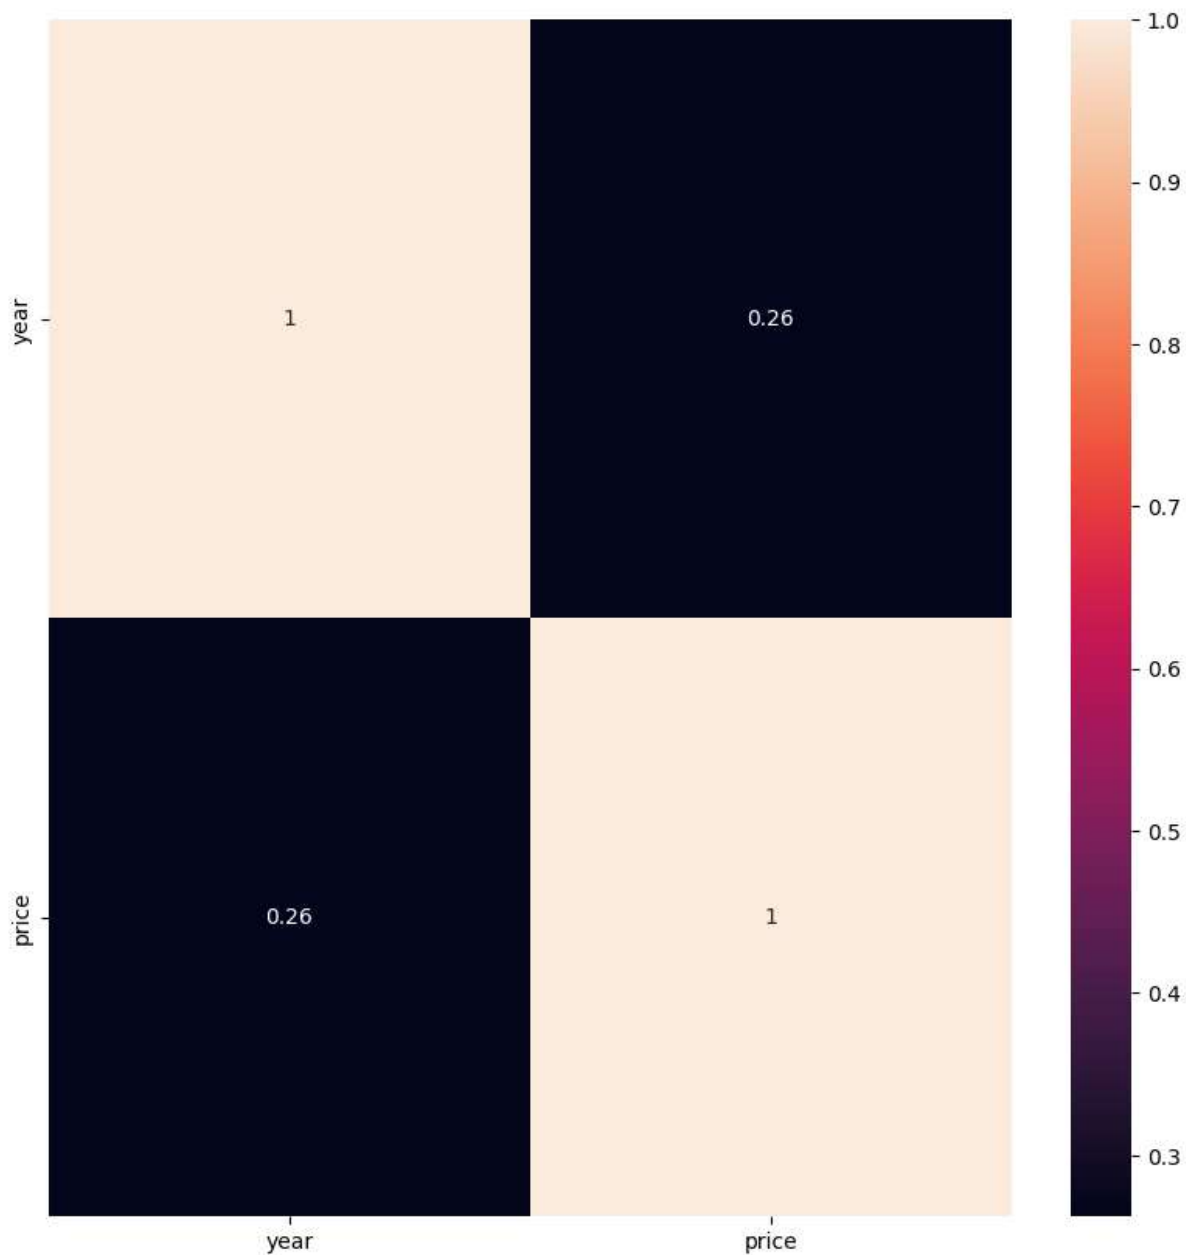


```
In [23]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(X_train,y_train)
y_pred=model.predict(X_test)
r2=r2_score(y_test,y_pred)
print("R2 score: ",r2)
```

R2 score: 0.17291193573932195

```
In [24]: plt.figure(figsize = (10, 10))  
sns.heatmap(df.corr(), annot = True)
```

Out[24]: <Axes: >



```
In [25]: import pandas as pd  
import numpy as np  
from sklearn import preprocessing  
import matplotlib.pyplot as plt  
import seaborn as sns  
sns.set(style="white")  
sns.set(style="whitegrid", color_codes=True)  
import warnings  
warnings.simplefilter(action='ignore')
```



```
In [27]: db=pd.read_csv(r"C:\Users\DELL\Downloads\used_cars_data.csv")
db
```

Out[27]:

	S.No.	Name	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	Owner_T
0	0	Maruti Wagon R LXI CNG	Mumbai	2010	72000	CNG	Manual	f
1	1	Hyundai Creta 1.6 CRDi SX Option	Pune	2015	41000	Diesel	Manual	f
2	2	Honda Jazz V	Chennai	2011	46000	Petrol	Manual	f
3	3	Maruti Ertiga VDI	Chennai	2012	87000	Diesel	Manual	f
4	4	Audi A4 New 2.0 TDI Multitronic	Coimbatore	2013	40670	Diesel	Automatic	Sec
...
7248	7248	Volkswagen Vento Diesel Trendline	Hyderabad	2011	89411	Diesel	Manual	f
7249	7249	Volkswagen Polo GT TSI	Mumbai	2015	59000	Petrol	Automatic	f
7250	7250	Nissan Micra Diesel XV	Kolkata	2012	28000	Diesel	Manual	f
7251	7251	Volkswagen Polo GT TSI	Pune	2013	52262	Petrol	Automatic	T
7252	7252	Mercedes-Benz E-Class 2009-2013 E 220 CDI Avan...	Kochi	2014	72443	Diesel	Automatic	f

7253 rows × 14 columns



In [32]: db.head()

Out[32]:

	S.No.	Name	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	Owner_Type
0	0	Maruti Wagon R LXI CNG	Mumbai	2010	72000	CNG	Manual	First
1	1	Hyundai Creta 1.6 CRDi SX Option	Pune	2015	41000	Diesel	Manual	First
2	2	Honda Jazz V	Chennai	2011	46000	Petrol	Manual	First
3	3	Maruti Ertiga VDI	Chennai	2012	87000	Diesel	Manual	First
4	4	Audi A4 New 2.0 TDI Multitronic	Coimbatore	2013	40670	Diesel	Automatic	Second

In [33]: db.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7253 entries, 0 to 7252
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   S.No.                 7253 non-null  int64
1   Name                  7253 non-null  object
2   Location              7253 non-null  object
3   Year                  7253 non-null  int64
4   Kilometers_Driven     7253 non-null  int64
5   Fuel_Type             7253 non-null  object
6   Transmission          7253 non-null  object
7   Owner_Type            7253 non-null  object
8   Mileage               7251 non-null  object
9   Engine                7207 non-null  object
10  Power                 7207 non-null  object
11  Seats                 7200 non-null  float64
12  New_Price             1006 non-null  object
13  Price                 6019 non-null  float64
dtypes: float64(2), int64(3), object(9)
memory usage: 793.4+ KB
```

In [34]: `db.describe()`

Out[34]:

	S.No.	Year	Kilometers_Driven	Seats	Price
count	7253.000000	7253.000000	7.253000e+03	7200.000000	6019.000000
mean	3626.000000	2013.365366	5.869906e+04	5.279722	9.479468
std	2093.905084	3.254421	8.442772e+04	0.811660	11.187917
min	0.000000	1996.000000	1.710000e+02	0.000000	0.440000
25%	1813.000000	2011.000000	3.400000e+04	5.000000	3.500000
50%	3626.000000	2014.000000	5.341600e+04	5.000000	5.640000
75%	5439.000000	2016.000000	7.300000e+04	5.000000	9.950000
max	7252.000000	2019.000000	6.500000e+06	10.000000	160.000000

In [35]: `db.shape`

Out[35]: (7253, 14)

In [36]: `db.isnull().sum()`

Out[36]:


S.No.	0
Name	0
Location	0
Year	0
Kilometers_Driven	0
Fuel_Type	0
Transmission	0
Owner_Type	0
Mileage	2
Engine	46
Power	46
Seats	53
New_Price	6247
Price	1234

dtype: int64

In [37]: `db.duplicated().any()`

Out[37]: False

```
In [38]: ax=db['Seats'].hist(bins=10,density=True,stacked=True,color='blue',alpha=0.6)
db['Seats'].plot(kind='density',color='violet')
ax.set(xlabel='Seats')
plt.xlim(-0,25)
plt.show()
```



```
In [39]: print(db["Seats"].mean(skipna=True))
print(db["Seats"].median(skipna=True))
```


```
5.279722222222222
5.0
```

```
In [40]: print(db["New_Price"].isnull().sum()/db.shape[0])
print(db["Price"].isnull().sum()/db.shape[0])
print(db["Mileage"].isnull().sum()/db.shape[0])
print(db["Engine"].isnull().sum()/db.shape[0])
print(db["Power"].isnull().sum()/db.shape[0])
```

```
0.8612987729215497
0.1701364952433476
0.0002757479663587481
0.006342203226251206
0.006342203226251206
```

```
In [41]: print(db['Engine'].value_counts())
sns.countplot(x='Engine',data=db,palette='Set3')
plt.xlim(-0,45)
plt.show()
```

```
Engine
1197 CC    732
1248 CC    610
1498 CC    370
998 CC     309
1198 CC    281
...
1489 CC     1
1422 CC     1
2706 CC     1
1978 CC     1
1389 CC     1
Name: count, Length: 150, dtype: int64
```



```
In [42]: data=db.copy()
data['Seats'].fillna(db['Seats'].median(skipna=True),inplace=True)
data.drop('New_Price',axis=1,inplace=True)
data['Price'].fillna(db['Price'].median(skipna=True),inplace=True)
data['Mileage'].fillna(db['Mileage'].value_counts().idxmax(),inplace=True)
data.drop('Engine',axis=1,inplace=True)
data.drop('Power',axis=1,inplace=True)
```


```
In [43]: data.isnull().any()
```

```
Out[43]: S.No.                False
Name                False
Location           False
Year               False
Kilometers_Driven  False
Fuel_Type          False
Transmission       False
Owner_Type         False
Mileage            False
Seats              False
Price              False
dtype: bool
```


```
In [44]: data.head()
```

```
Out[44]:
```

	S.No.	Name	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	Owner_Type
0	0	Maruti Wagon R LXI CNG	Mumbai	2010	72000	CNG	Manual	First
1	1	Hyundai Creta 1.6 CRDi SX Option	Pune	2015	41000	Diesel	Manual	First
2	2	Honda Jazz V	Chennai	2011	46000	Petrol	Manual	First
3	3	Maruti Ertiga VDI	Chennai	2012	87000	Diesel	Manual	First
4	4	Audi A4 New 2.0 TDI Multitronic	Coimbatore	2013	40670	Diesel	Automatic	Second



```
In [45]: plt.figure(figsize=(15,8))
ax=db["Price"].hist(bins=10,density=True,stacked=True,color='blue',alpha=0.6)
db["Price"].plot(kind='density',color='teal')
ax.legend(['Mileage','Seats'])
ax.set(xlabel='Price')
plt.xlim(-9,80)
plt.show()
```




```
In [46]: training=pd.get_dummies(data,columns=["S.No."])
final_train=training
final_train.head()
```

Out[46]:

	Name	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	Owner_Type	Mileage
0	Maruti Wagon R LXI CNG	Mumbai	2010	72000	CNG	Manual	First	26.6 km/kg
1	Hyundai Creta 1.6 CRDi SX Option	Pune	2015	41000	Diesel	Manual	First	19.6 km/kg
2	Honda Jazz V	Chennai	2011	46000	Petrol	Manual	First	18.2 km/kg
3	Maruti Ertiga VDI	Chennai	2012	87000	Diesel	Manual	First	20.7 km/kg
4	Audi A4 New 2.0 TDI Multitronic	Coimbatore	2013	40670	Diesel	Automatic	Second	15.2 km/kg


5 rows × 7263 columns



```
In [47]: sns.barplot(x='Price',y='Year',data=final_train,color='mediumturquoise')  
plt.show()
```




```
In [48]: import seaborn as sns
import matplotlib.pyplot as plt
sns.barplot(x='Year',y='Seats',data=db,color='aquamarine')
plt.show()
```



In []: